

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler,
RobustScaler, Normalizer, LabelEncoder, OneHotEncoder
from sklearn.impute import SimpleImputer, KNNImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier, AdaBoostClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix, classification_report
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score
from sklearn.metrics import roc_auc_score, roc_curve

```

```

df=pd.read_csv(r"C:\Users\arjun\Downloads\BankChurners (1).csv")
df.head()

```

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count
0	768805383	Existing Customer	45	M	3
1	818770008	Existing Customer	49	F	5
2	713982108	Existing Customer	51	M	3
3	769911858	Existing Customer	40	F	4
4	709106358	Existing Customer	40	M	3

	Education_Level	Marital_Status	Income_Category	Card_Category
0	High School	Married	\$60K - \$80K	Blue
1	Graduate	Single	Less than \$40K	Blue
2	Graduate	Married	\$80K - \$120K	Blue
3	High School	Unknown	Less than \$40K	Blue
4	Uneducated	Married	\$60K - \$80K	Blue

	Months_on_book	...	Months_Inactive_12_mon	Contacts_Count_12_mon
0	39	...	1	3
1	44	...	1	2

2	36	...	1	0
3	34	...	4	1
4	21	...	1	0

	Credit_Limit	Total_Revolving_Bal	Avg_Open_To_Buy
Total_Amt_Chng_Q4_Q1 \			
0	12691.0	777	11914.0
1.335			
1	8256.0	864	7392.0
1.541			
2	3418.0	0	3418.0
2.594			
3	3313.0	2517	796.0
1.405			
4	4716.0	0	4716.0
2.175			

	Total_Trans_Amt	Total_Trans_Ct	Total_Ct_Chng_Q4_Q1
Avg_Utilization_Ratio			
0	1144	42	1.625
0.061			
1	1291	33	3.714
0.105			
2	1887	20	2.333
0.000			
3	1171	20	2.333
0.760			
4	816	28	2.500
0.000			

[5 rows x 21 columns]

```
df.isna().sum()
```

CLIENTNUM	0
Attrition_Flag	0
Customer_Age	0
Gender	0
Dependent_count	0
Education_Level	0
Marital_Status	0
Income_Category	0
Card_Category	0
Months_on_book	0
Total_Relationship_Count	0
Months_Inactive_12_mon	0
Contacts_Count_12_mon	0

```
Credit_Limit          0
Total_Revolving_Bal   0
Avg_Open_To_Buy       0
Total_Amt_Chng_Q4_Q1  0
Total_Trans_Amt       0
Total_Trans_Ct        0
Total_Ct_Chng_Q4_Q1   0
Avg_Utilization_Ratio 0
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10127 entries, 0 to 10126
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	CLIENTNUM	10127 non-null	int64
1	Attrition_Flag	10127 non-null	object
2	Customer_Age	10127 non-null	int64
3	Gender	10127 non-null	object
4	Dependent_count	10127 non-null	int64
5	Education_Level	10127 non-null	object
6	Marital_Status	10127 non-null	object
7	Income_Category	10127 non-null	object
8	Card_Category	10127 non-null	object
9	Months_on_book	10127 non-null	int64
10	Total_Relationship_Count	10127 non-null	int64
11	Months_Inactive_12_mon	10127 non-null	int64
12	Contacts_Count_12_mon	10127 non-null	int64
13	Credit_Limit	10127 non-null	float64
14	Total_Revolving_Bal	10127 non-null	int64
15	Avg_Open_To_Buy	10127 non-null	float64
16	Total_Amt_Chng_Q4_Q1	10127 non-null	float64
17	Total_Trans_Amt	10127 non-null	int64
18	Total_Trans_Ct	10127 non-null	int64
19	Total_Ct_Chng_Q4_Q1	10127 non-null	float64
20	Avg_Utilization_Ratio	10127 non-null	float64

```
dtypes: float64(5), int64(10), object(6)
```

```
memory usage: 1.6+ MB
```

```
df['Attrition_Flag'].value_counts()
```

```
Attrition_Flag
Existing Customer    8500
Attrited Customer    1627
Name: count, dtype: int64
```

```
df['Attrition_Flag'] = df['Attrition_Flag'].map({'Existing Customer':
0, 'Attrited Customer': 1})
```

```
df['Gender'].value_counts()
```

Gender

F 5358

M 4769

Name: count, dtype: int64

```
df['Gender'] = df['Gender'].map({'F': 0, 'M': 1})
```

```
df = pd.get_dummies(df, columns=['Education_Level'], dtype=int)
```

```
df.head()
```

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	\
0	768805383	0	45	1	3	
1	818770008	0	49	0	5	
2	713982108	0	51	1	3	
3	769911858	0	40	0	4	
4	709106358	0	40	1	3	

	Marital_Status	Income_Category	Card_Category	Months_on_book	\
0	Married	\$60K - \$80K	Blue	39	
1	Single	Less than \$40K	Blue	44	
2	Married	\$80K - \$120K	Blue	36	
3	Unknown	Less than \$40K	Blue	34	
4	Married	\$60K - \$80K	Blue	21	

	Total_Relationship_Count	...	Total_Trans_Ct	Total_Ct_Chng_Q4_Q1	\
0	5	...	42	1.625	
1	6	...	33	3.714	
2	4	...	20	2.333	
3	3	...	20	2.333	
4	5	...	28	2.500	

	Avg_Utilization_Ratio	Education_Level_College	Education_Level_Doctorate	\
0	0.061		0	
0				
1	0.105		0	
0				
2	0.000		0	
0				
3	0.760		0	
0				

```

4          0.000          0
0

Education_Level_Graduate Education_Level_High School \
0          0          1
1          1          0
2          1          0
3          0          1
4          0          0

Education_Level_Post-Graduate Education_Level_Uneducated \
0          0          0
1          0          0
2          0          0
3          0          0
4          0          1

Education_Level_Unknown
0          0
1          0
2          0
3          0
4          0

[5 rows x 27 columns]

df = pd.get_dummies(df, columns=['Marital_Status'], dtype=int)
df.head()

CLIENTNUM  Attrition_Flag  Customer_Age  Gender  Dependent_count \
0  768805383          0          45          1          3
1  818770008          0          49          0          5
2  713982108          0          51          1          3
3  769911858          0          40          0          4
4  709106358          0          40          1          3

Income_Category  Card_Category  Months_on_book
Total_Relationship_Count \
0    $60K - $80K          Blue          39
5
1  Less than $40K          Blue          44
6
2    $80K - $120K          Blue          36
4
3  Less than $40K          Blue          34
3
4    $60K - $80K          Blue          21
5

Months_Inactive_12_mon  ...  Education_Level_Doctorate \

```

0	1	...	0
1	1	...	0
2	1	...	0
3	4	...	0
4	1	...	0

	Education_Level_Graduate	Education_Level_High School \
0	0	1
1	1	0
2	1	0
3	0	1
4	0	0

	Education_Level_Post-Graduate	Education_Level_Uneducated \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	1

	Education_Level_Unknown	Marital_Status_Divorced
Marital_Status_Married \		
0	0	0
1		
1	0	0
0		
2	0	0
1		
3	0	0
0		
4	0	0
1		

	Marital_Status_Single	Marital_Status_Unknown
0	0	0
1	1	0
2	0	0
3	0	1
4	0	0

[5 rows x 30 columns]

```
income_mapping = {
    'Unknown': 0,
    'Less than $40K': 1,
    '$40K - $60K': 2,
    '$60K - $80K': 3,
    '$80K - $120K': 4,
    '$120K +': 5
}
```

```
df['Income_Category'] = df['Income_Category'].map(income_mapping)
df.head()
```

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	\
0	768805383	0	45	1	3	
1	818770008	0	49	0	5	
2	713982108	0	51	1	3	
3	769911858	0	40	0	4	
4	709106358	0	40	1	3	

	Income_Category	Card_Category	Months_on_book	Total_Relationship_Count	\
0		3	Blue	39	
5					
1		1	Blue	44	
6					
2		4	Blue	36	
4					
3		1	Blue	34	
3					
4		3	Blue	21	
5					

	Months_Inactive_12_mon	...	Education_Level_Doctorate	\
0	1	...	0	
1	1	...	0	
2	1	...	0	
3	4	...	0	
4	1	...	0	

	Education_Level_Graduate	Education_Level_High School	\
0	0	1	
1	1	0	
2	1	0	
3	0	1	
4	0	0	

	Education_Level_Post-Graduate	Education_Level_Uneducated	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	1	

	Education_Level_Unknown	Marital_Status_Divorced	Marital_Status_Married	\
0		0		0
1				

1	0	0
0		
2	0	0
1		
3	0	0
0		
4	0	0
1		

	Marital_Status_Single	Marital_Status_Unknown
0	0	0
1	1	0
2	0	0
3	0	1
4	0	0

[5 rows x 30 columns]

```
df = pd.get_dummies(df, columns=['Card_Category'], dtype=int)
df.head()
```

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	\
0	768805383	0	45	1	3	
1	818770008	0	49	0	5	
2	713982108	0	51	1	3	
3	769911858	0	40	0	4	
4	709106358	0	40	1	3	

	Income_Category	Months_on_book	Total_Relationship_Count	\
0	3	39	5	
1	1	44	6	
2	4	36	4	
3	1	34	3	
4	3	21	5	

	Months_Inactive_12_mon	Contacts_Count_12_mon	...	\
0	1	3	...	
1	1	2	...	
2	1	0	...	
3	4	1	...	
4	1	0	...	

	Education_Level_Uneducated	Education_Level_Unknown	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	1	0	

	Marital_Status_Divorced	Marital_Status_Married
--	-------------------------	------------------------

Marital_Status_Single \		
0	0	1
0		
1	0	0
1		
2	0	1
0		
3	0	0
0		
4	0	1
0		

Marital_Status_Unknown	Card_Category_Blue	Card_Category_Gold \
0	0	1 0
1	0	1 0
2	0	1 0
3	1	1 0
4	0	1 0

Card_Category_Platinum	Card_Category_Silver
0	0
1	0
2	0
3	0
4	0

[5 rows x 33 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10127 entries, 0 to 10126

Data columns (total 33 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	CLIENTNUM	10127 non-null	int64
1	Attrition_Flag	10127 non-null	int64
2	Customer_Age	10127 non-null	int64
3	Gender	10127 non-null	int64
4	Dependent_count	10127 non-null	int64
5	Income_Category	10127 non-null	int64
6	Months_on_book	10127 non-null	int64
7	Total_Relationship_Count	10127 non-null	int64
8	Months_Inactive_12_mon	10127 non-null	int64
9	Contacts_Count_12_mon	10127 non-null	int64
10	Credit_Limit	10127 non-null	float64
11	Total_Revolving_Bal	10127 non-null	int64
12	Avg_Open_To_Buy	10127 non-null	float64
13	Total_Amt_Chng_Q4_Q1	10127 non-null	float64
14	Total_Trans_Amt	10127 non-null	int64

```

15 Total_Trans_Ct                10127 non-null   int64
16 Total_Ct_Chng_Q4_Q1          10127 non-null   float64
17 Avg_Utilization_Ratio        10127 non-null   float64
18 Education_Level_College       10127 non-null   int32
19 Education_Level_Doctorate     10127 non-null   int32
20 Education_Level_Graduate      10127 non-null   int32
21 Education_Level_High School   10127 non-null   int32
22 Education_Level_Post-Graduate 10127 non-null   int32
23 Education_Level_Uneducated    10127 non-null   int32
24 Education_Level_Unknown       10127 non-null   int32
25 Marital_Status_Divorced       10127 non-null   int32
26 Marital_Status_Married        10127 non-null   int32
27 Marital_Status_Single         10127 non-null   int32
28 Marital_Status_Unknown        10127 non-null   int32
29 Card_Category_Blue            10127 non-null   int32
30 Card_Category_Gold            10127 non-null   int32
31 Card_Category_Platinum        10127 non-null   int32
32 Card_Category_Silver          10127 non-null   int32
dtypes: float64(5), int32(15), int64(13)
memory usage: 2.0 MB

df.drop(columns=["CLIENTNUM"], inplace=True)

from scipy.stats import skew

# List of columns to check for skewness
skew_columns = [
    "Credit_Limit", "Total_Revolving_Bal", "Avg_Open_To_Buy",
    "Total_Amt_Chng_Q4_Q1",
    "Total_Trans_Amt", "Total_Trans_Ct", "Total_Ct_Chng_Q4_Q1",
    "Avg_Utilization_Ratio"
]

# Calculate skewness for each column
skewness_values = {col: skew(df[col], nan_policy='omit') for col in
skew_columns}

# Display skewness values
for col, skew_val in skewness_values.items():
    print(f"Skewness of {col}: {skew_val:.2f}")

Skewness of Credit_Limit: 1.67
Skewness of Total_Revolving_Bal: -0.15
Skewness of Avg_Open_To_Buy: 1.66
Skewness of Total_Amt_Chng_Q4_Q1: 1.73
Skewness of Total_Trans_Amt: 2.04
Skewness of Total_Trans_Ct: 0.15
Skewness of Total_Ct_Chng_Q4_Q1: 2.06
Skewness of Avg_Utilization_Ratio: 0.72

```

```

"""import numpy as np

# Apply log transformation for highly skewed columns
columns_to_transform = ["Credit_Limit", "Avg_Open_To_Buy",
"Total_Amt_Chng_Q4_Q1", "Total_Trans_Amt", "Total_Ct_Chng_Q4_Q1"]

for col in columns_to_transform:
    df[col] = np.log1p(df[col]) # log1p to avoid log(0) issues

print("Transformation applied to highly skewed columns!")"""

'import numpy as np\n\n# Apply log transformation for highly skewed
columns\ncolumns_to_transform = ["Credit_Limit", "Avg_Open_To_Buy",
"Total_Amt_Chng_Q4_Q1", "Total_Trans_Amt", "Total_Ct_Chng_Q4_Q1"]\n\n
nfor col in columns_to_transform:\n    df[col] = np.log1p(df[col]) #
log1p to avoid log(0) issues\n\nprint("Transformation applied to
highly skewed columns!")'

"""from sklearn.preprocessing import RobustScaler

scaler = RobustScaler()
df[['Credit_Limit', 'Total_Revolving_Bal', 'Avg_Open_To_Buy',
'Total_Trans_Amt']] = scaler.fit_transform(df[['Credit_Limit',
'Total_Revolving_Bal', 'Avg_Open_To_Buy', 'Total_Trans_Amt']])"""

"from sklearn.preprocessing import RobustScaler\n\nscaler =
RobustScaler()\nndf[['Credit_Limit', 'Total_Revolving_Bal',
'Avg_Open_To_Buy', 'Total_Trans_Amt']] =
scaler.fit_transform(df[['Credit_Limit', 'Total_Revolving_Bal',
'Avg_Open_To_Buy', 'Total_Trans_Amt']])"

"""import numpy as np

# List of columns to check for outliers
outlier_columns = [
    "Credit_Limit", "Total_Revolving_Bal", "Avg_Open_To_Buy",
"Total_Amt_Chng_Q4_Q1",
    "Total_Trans_Amt", "Total_Trans_Ct", "Total_Ct_Chng_Q4_Q1",
"Avg_Utilization_Ratio"
]

# Function to remove outliers using IQR
def remove_outliers(df, columns):
    for col in columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]

```

```

    return df

# Apply outlier removal
df1 = remove_outliers(df, outlier_columns)

# Display the new shape of the dataset
print(f"Original shape: {df.shape}")
print(f"New shape after outlier removal: {df1.shape}")"""

'import numpy as np\n\n# List of columns to check for outliers\n
noutlier_columns = [\n    "Credit_Limit", "Total_Revolving_Bal",
"Avg_Open_To_Buy", "Total_Amt_Chng_Q4_Q1",\n    "Total_Trans_Amt",
"Total_Trans_Ct", "Total_Ct_Chng_Q4_Q1", "Avg_Utilization_Ratio"\n]\n\n
# Function to remove outliers using IQR\ndef remove_outliers(df,
columns):\n    for col in columns:\n        Q1 =
df[col].quantile(0.25)\n        Q3 = df[col].quantile(0.75)\n
IQR = Q3 - Q1\n        lower_bound = Q1 - 1.5 * IQR\n
upper_bound = Q3 + 1.5 * IQR\n        df = df[(df[col] >= lower_bound)
& (df[col] <= upper_bound)]\n    return df\n\n# Apply outlier removal\n
ndf1 = remove_outliers(df, outlier_columns)\n\n# Display the new shape
of the dataset\nprint(f"Original shape: {df.shape}")\nprint(f"New
shape after outlier removal: {df1.shape}")'

import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd






# Select numeric columns
numeric_cols = df.select_dtypes(include=['int64', 'float64',
'int32']).columns

# Create subplots
fig, axes = plt.subplots(nrows=len(numeric_cols), ncols=1,
figsize=(10, len(numeric_cols) * 3))

# Loop through each column and plot in a separate subplot
for i, column in enumerate(numeric_cols):
    sns.kdeplot(df[column], ax=axes[i], fill=True)
    axes[i].set_title(f"Distribution of {column}")

# Adjust layout
plt.tight_layout()
plt.show()

```








```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Select numeric columns
numeric_cols = df.select_dtypes(include=['int64', 'float64',
'int32']).columns

# Create subplots
fig, axes = plt.subplots(nrows=len(numeric_cols), ncols=1,
figsize=(10, len(numeric_cols) * 3))

# Loop through each column and plot in a separate subplot
for i, column in enumerate(numeric_cols):
    sns.boxplot(df[column], ax=axes[i], fill=True)
    axes[i].set_title(f"Distribution of {column}")

# Adjust layout
plt.tight_layout()
plt.show()
```



```
df.head()
```

	Attrition_Flag	Customer_Age	Gender	Dependent_count
Income_Category \				
0	0	45	1	3
3				
1	0	49	0	5
1				
2	0	51	1	3
4				
3	0	40	0	4
1				
4	0	40	1	3
3				

	Months_on_book	Total_Relationship_Count	Months_Inactive_12_mon	\
0	39	5	1	
1	44	6	1	
2	36	4	1	
3	34	3	4	
4	21	5	1	

	Contacts_Count_12_mon	Credit_Limit	...
Education_Level_Uneducated \			
0	3	12691.0	...
0			
1	2	8256.0	...
0			
2	0	3418.0	...
0			
3	1	3313.0	...
0			
4	0	4716.0	...
1			

	Education_Level_Unknown	Marital_Status_Divorced
Marital_Status_Married \		
0	0	0
1		
1	0	0
0		
2	0	0
1		
3	0	0
0		
4	0	0
1		

	Marital_Status_Single	Marital_Status_Unknown	
Card_Category_Blue \			
0	0	0	1
1	1	0	1
2	0	0	1
3	0	1	1
4	0	0	1

	Card_Category_Gold	Card_Category_Platinum	Card_Category_Silver
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

[5 rows x 32 columns]

```
correlation = df.corr()["Attrition_Flag"].sort_values(ascending=False)
print(correlation)
```

Attrition_Flag	1.000000
Contacts_Count_12_mon	0.204491
Months_Inactive_12_mon	0.152449
Education_Level_Doctorate	0.029386
Marital_Status_Single	0.019037
Dependent_count	0.018991
Customer_Age	0.018203
Months_on_book	0.013687
Education_Level_Post-Graduate	0.011127
Card_Category_Platinum	0.010823
Education_Level_Unknown	0.009005
Marital_Status_Unknown	0.008904
Card_Category_Gold	0.005973
Card_Category_Blue	0.003216
Marital_Status_Divorced	0.000850
Avg_Open_To_Buy	-0.000285
Education_Level_Uneducated	-0.001444
Education_Level_College	-0.007840
Card_Category_Silver	-0.008467
Education_Level_Graduate	-0.009046
Education_Level_High School	-0.011730
Income_Category	-0.013577
Marital_Status_Married	-0.023735
Credit_Limit	-0.023873

```

Gender -0.037272
Total_Amt_Chng_Q4_Q1 -0.131063
Total_Relationship_Count -0.150005
Total_Trans_Amt -0.168598
Avg_Utilization_Ratio -0.178410
Total_Revolving_Bal -0.263053
Total_Ct_Chng_Q4_Q1 -0.290054
Total_Trans_Ct -0.371403
Name: Attrition_Flag, dtype: float64

```

```
from scipy.stats import skew
```

```
# List of columns to check for skewness
```

```
skew_columns = [
    "Total_Revolving_Bal", "Total_Ct_Chng_Q4_Q1",
    "Total_Trans_Ct",
    "Contacts_Count_12_mon", 'Months_Inactive_12_mon'
]
```

```
# Calculate skewness for each column
```

```
skewness_values = {col: skew(df[col], nan_policy='omit') for col in
skew_columns}
```

```
# Display skewness values
```

```
for col, skew_val in skewness_values.items():
    print(f"Skewness of {col}: {skew_val:.2f}")
```

```

Skewness of Total_Revolving_Bal: -0.15
Skewness of Total_Ct_Chng_Q4_Q1: 2.06
Skewness of Total_Trans_Ct: 0.15
Skewness of Contacts_Count_12_mon: 0.01
Skewness of Months_Inactive_12_mon: 0.63

```

```
df.describe()
```

	Attrition_Flag	Customer_Age	Gender	Dependent_count \
count	10127.000000	10127.000000	10127.000000	10127.000000
mean	0.160660	46.325960	0.470919	2.346203
std	0.367235	8.016814	0.499178	1.298908
min	0.000000	26.000000	0.000000	0.000000
25%	0.000000	41.000000	0.000000	1.000000
50%	0.000000	46.000000	0.000000	2.000000
75%	0.000000	52.000000	1.000000	3.000000
max	1.000000	73.000000	1.000000	5.000000

	Income_Category	Months_on_book	Total_Relationship_Count \
count	10127.000000	10127.000000	10127.000000
mean	2.085711	35.928409	3.812580
std	1.474639	7.986416	1.554408
min	0.000000	13.000000	1.000000
25%	1.000000	31.000000	3.000000

50%	2.000000	36.000000	4.000000
75%	3.000000	40.000000	5.000000
max	5.000000	56.000000	6.000000

	Months_Inactive_12_mon	Contacts_Count_12_mon
Credit_Limit \		
count	10127.000000	10127.000000
10127.000000 ...		
mean	2.341167	2.455317
8631.953698 ...		
std	1.010622	1.106225
9088.776650 ...		
min	0.000000	0.000000
1438.300000 ...		
25%	2.000000	2.000000
2555.000000 ...		
50%	2.000000	2.000000
4549.000000 ...		
75%	3.000000	3.000000
11067.500000 ...		
max	6.000000	6.000000
34516.000000 ...		

	Education_Level_Uneducated	Education_Level_Unknown \
count	10127.000000	10127.000000
mean	0.146835	0.149995
std	0.353959	0.357084
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	Marital_Status_Divorced	Marital_Status_Married
Marital_Status_Single \		
count	10127.000000	10127.000000
10127.000000		
mean	0.073862	0.462822
0.389355		
std	0.261559	0.498641
0.487628		
min	0.000000	0.000000
0.000000		
25%	0.000000	0.000000
0.000000		
50%	0.000000	0.000000
0.000000		
75%	0.000000	1.000000
1.000000		
max	1.000000	1.000000

1.000000

	Marital_Status_Unknown	Card_Category_Blue	Card_Category_Gold
\count	10127.000000	10127.000000	10127.000000
mean	0.073961	0.931767	0.011455
std	0.261720	0.252159	0.106416
min	0.000000	0.000000	0.000000
25%	0.000000	1.000000	0.000000
50%	0.000000	1.000000	0.000000
75%	0.000000	1.000000	0.000000
max	1.000000	1.000000	1.000000

	Card_Category_Platinum	Card_Category_Silver
count	10127.000000	10127.000000
mean	0.001975	0.054804
std	0.044398	0.227608
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

[8 rows x 32 columns]

model

```
"""selected_features = ['Total_Amt_Chng_Q4_Q1',
'Total_Relationship_Count', 'Total_Trans_Amt',
'Avg_Utilization_Ratio', 'Total_Revolving_Bal',
'Total_Ct_Chng_Q4_Q1',
'Total_Trans_Ct', 'Contacts_Count_12_mon',
'Months_Inactive_12_mon']"""
selected_features = [ 'Total_Relationship_Count',
'Total_Revolving_Bal', 'Total_Ct_Chng_Q4_Q1',
'Total_Trans_Ct', 'Contacts_Count_12_mon',
'Months_Inactive_12_mon']
X = df[selected_features]
y = df["Attrition_Flag"] # Target variable

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```

from imblearn.over_sampling import SMOTE
smote = SMOTE(sampling_strategy=0.5, random_state=42) # Only 50%
oversampling
X_train_resampled, y_train_resampled = smote.fit_resample(X_train,
y_train)

```

Train logistic regression model on balanced data

```

model = LogisticRegression()
model.fit(X_train_resampled, y_train_resampled)

```

Make predictions

```

y_pred = model.predict(X_test)

```

Evaluation

```

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

```

Print results

```

print(f"Accuracy after SMOTE: {accuracy * 100:.2f}%")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)

```

Accuracy after SMOTE: 88.25%

Confusion Matrix:

```

[[1563  136]
 [ 102  225]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.92	0.93	1699
1	0.62	0.69	0.65	327
accuracy			0.88	2026
macro avg	0.78	0.80	0.79	2026
weighted avg	0.89	0.88	0.88	2026

C:\Users\arjun\AppData\Roaming\Python\Python310\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

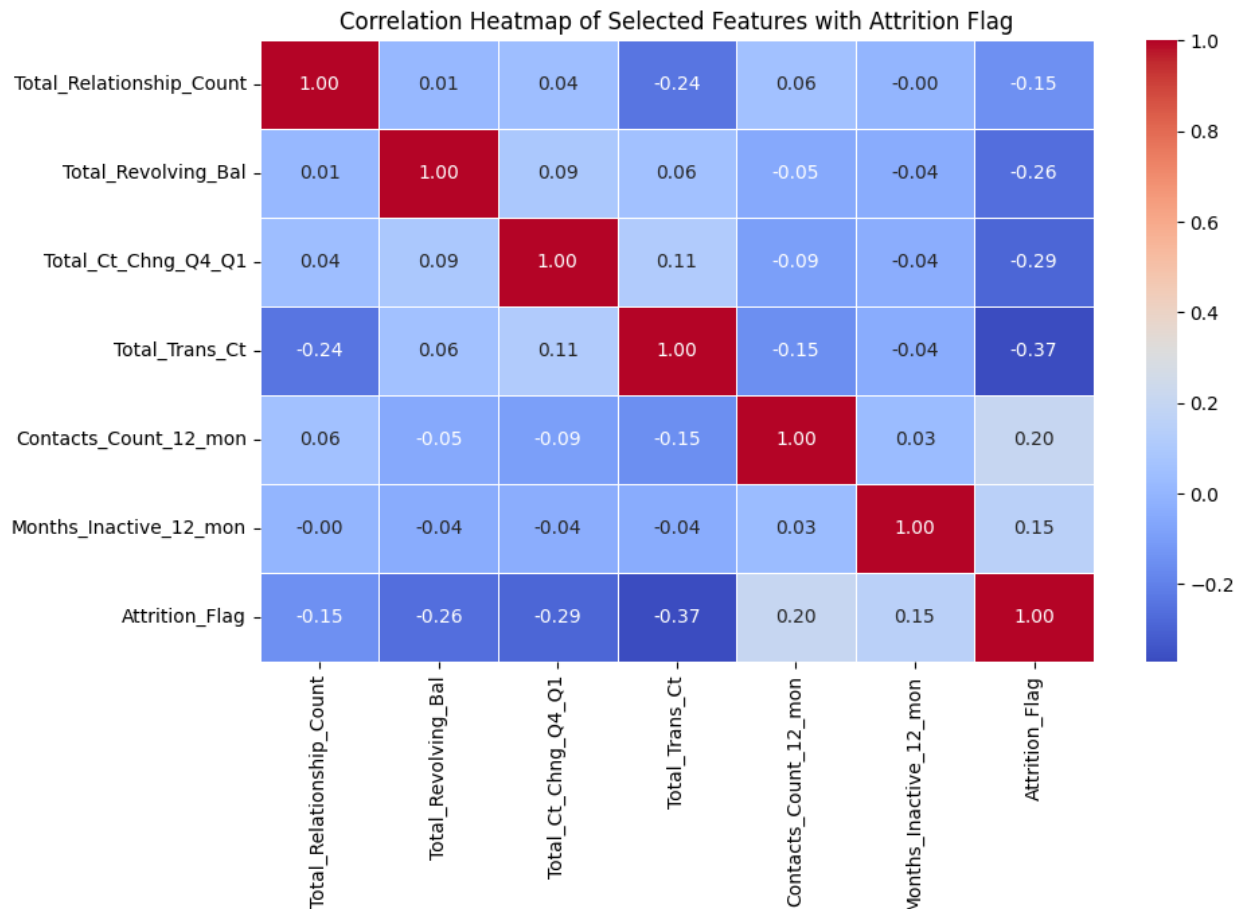
```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
    n_iter_i = _check_optimize_result(

selected_features = [ 'Total_Relationship_Count',
                      'Total_Revolving_Bal', 'Total_Ct_Chng_Q4_Q1',
                      'Total_Trans_Ct', 'Contacts_Count_12_mon',
                      'Months_Inactive_12_mon', 'Attrition_Flag']

# Filter dataset with selected features
df_selected = df[selected_features]

# Compute correlation matrix
corr_matrix = df_selected.corr()

# Plot heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=0.5)
plt.title("Correlation Heatmap of Selected Features with Attrition
Flag")
plt.show()
```




```
selected_features = ['Total_Amt_Chng_Q4_Q1',
'Total_Relationship_Count', 'Total_Trans_Amt',
'Avg_Utilization_Ratio', 'Total_Revolving_Bal',
'Total_Ct_Chng_Q4_Q1',
'Total_Trans_Ct', 'Contacts_Count_12_mon',
'Months_Inactive_12_mon', 'Attrition_Flag']

# Filter dataset with selected features
df_selected = df[selected_features]

# Compute correlation matrix
corr_matrix = df_selected.corr()

# Plot heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f",
linewidths=0.5)
plt.title("Correlation Heatmap of Selected Features with Attrition
Flag")
plt.show()
```



```
import numpy as np
from sklearn.preprocessing import MinMaxScaler

# Load your trained model and scaler
# model = your_trained_model
# scaler = your_fitted_scaler

# Selected features as per your specification
selected_features = [
    'Total_Relationship_Count',
    'Total_Revolving_Bal',
    'Total_Ct_Chng_Q4_Q1',
    'Total_Trans_Ct',
    'Contacts_Count_12_mon',
    'Months_Inactive_12_mon'
]

# Creating a single sample using the values you provided
new_sample = np.array([
    5,          # Total_Relationship_Count
    777,        # Total_Revolving_Bal
    1.625,      # Total_Ct_Chng_Q4_Q1
])
```



```

    42,      # Total_Trans_Ct
    3,      # Contacts_Count_12_mon
    1       # Months_Inactive_12_mon
11)

# Scale the features using your trained scaler
# scaled_sample = scaler.transform(new_sample)

# Make prediction
# prediction = model.predict(scaled_sample)
# probability = model.predict_proba(scaled_sample)[0][1] #
Probability of attrition (class 1)

# For demonstration (replace with actual prediction)
prediction = model.predict(new_sample) # Example prediction (0=stay,
1=attrite)

print("Customer data:", new_sample[0])
print("Prediction:", "Will attrite" if prediction == 1 else "Will
stay")

Customer data: [ 5.    777.    1.625  42.    3.    1. ]
Prediction: Will stay

```

C:\Users\arjun\AppData\Roaming\Python\Python310\site-packages\sklearn\
base.py:493: UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
 warnings.warn(

```

import numpy as np
from sklearn.preprocessing import MinMaxScaler

# Load your trained model and scaler
# model = your_trained_model
# scaler = your_fitted_scaler

# Selected features as per your specification
selected_features = [
    'Total_Relationship_Count',
    'Total_Revolving_Bal',
    'Total_Ct_Chng_Q4_Q1',
    'Total_Trans_Ct',
    'Contacts_Count_12_mon',
    'Months_Inactive_12_mon'
]

# Creating a single sample using the values you provided
new_sample = np.array([
    2,      # Total_Relationship_Count
    0,      # Total_Revolving_Bal

```

```

    0.6,    # Total_Ct_Chng_Q4_Q1
    16,    # Total_Trans_Ct
    3,     # Contacts_Count_12_mon
    3      # Months_Inactive_12_mon
11)

# Scale the features using your trained scaler
# scaled_sample = scaler.transform(new_sample)

# Make prediction
# prediction = model.predict(scaled_sample)
# probability = model.predict_proba(scaled_sample)[0][1]  #
Probability of attrition (class 1)

# For demonstration (replace with actual prediction)
prediction = model.predict(new_sample)  # Example prediction (0=stay,
1=attrite)

print("Customer data:", new_sample[0])
print("Prediction:", "Will attrite" if prediction == 1 else "Will
stay")

Customer data: [ 2.    0.    0.6 16.    3.    3. ]
Prediction: Will attrite

C:\Users\arjun\AppData\Roaming\Python\Python310\site-packages\sklearn\
base.py:493: UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
  warnings.warn(

```

KNN

```

"""from imblearn.over_sampling import SMOTE
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
from sklearn.preprocessing import RobustScaler # For scaling with
outliers

# Apply robust scaling that's less sensitive to outliers
robust_scaler = RobustScaler()
X_train_scaled = robust_scaler.fit_transform(X_train)
X_test_scaled = robust_scaler.transform(X_test)

# Apply SMOTE on scaled data
smote = SMOTE(sampling_strategy=0.5, random_state=42)
X_train_resampled, y_train_resampled =
smote.fit_resample(X_train_scaled, y_train)

```

```

# Train KNN model on scaled and balanced data
model = KNeighborsClassifier(n_neighbors=5)
model.fit(X_train_resampled, y_train_resampled)

# Make predictions on scaled test data
y_pred = model.predict(X_test_scaled)

# Evaluation
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

# Print results
print(f"Accuracy after robust scaling, SMOTE, with KNN: {accuracy * 100:.2f}%")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)"""

'from imblearn.over_sampling import SMOTE\nfrom sklearn.neighbors\nimport KNeighborsClassifier\nfrom sklearn.metrics import\naccuracy_score, confusion_matrix, classification_report\nfrom\nsklearn.preprocessing import RobustScaler # For scaling with\noutliers\n\n# Apply robust scaling that's less sensitive to outliers\n\nrobust_scaler = RobustScaler()\nX_train_scaled =\nrobust_scaler.fit_transform(X_train)\nX_test_scaled =\nrobust_scaler.transform(X_test)\n\n# Apply SMOTE on scaled data\nsmote\n= SMOTE(sampling_strategy=0.5, random_state=42)\nX_train_resampled,\ny_train_resampled = smote.fit_resample(X_train_scaled, y_train)\n\n# Train KNN model on scaled and balanced data\nmodel =\nKNeighborsClassifier(n_neighbors=5)\nmodel.fit(X_train_resampled,\ny_train_resampled)\n\n# Make predictions on scaled test data\nny_pred =\nmodel.predict(X_test_scaled)\n\n# Evaluation\nnaccuracy =\naccuracy_score(y_test, y_pred)\n\nconf_matrix = confusion_matrix(y_test,\ny_pred)\n\nnclass_report = classification_report(y_test, y_pred)\n\n# Print results\n\nprint(f"Accuracy after robust scaling, SMOTE, with KNN:\n{accuracy * 100:.2f}%")\n\nprint("Confusion Matrix:")\n\nprint(conf_matrix)\n\nprint("Classification Report:")\n\nprint(class_report)'\n\nimport numpy as np\nimport matplotlib.pyplot as plt\nfrom sklearn.model_selection import cross_val_score\nfrom imblearn.over_sampling import SMOTE\nfrom sklearn.neighbors import KNeighborsClassifier\nfrom sklearn.preprocessing import RobustScaler\n\n# Step 1: Apply RobustScaler to handle outliers\nrobust_scaler = RobustScaler()

```

```

X_train_scaled = robust_scaler.fit_transform(X_train)
X_test_scaled = robust_scaler.transform(X_test)

# Step 2: Apply SMOTE on the scaled training data
smote = SMOTE(sampling_strategy=0.5, random_state=42)
X_train_resampled, y_train_resampled =
smote.fit_resample(X_train_scaled, y_train)

# Step 3: Cross-validation to find the best k
k_values = range(1, 21) # Testing k from 1 to 20
cv_scores = []

for k in k_values:
    model = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(model, X_train_resampled,
y_train_resampled, cv=5, scoring='accuracy') # 5-fold CV
    cv_scores.append(scores.mean())

# Find the best k
best_k = k_values[np.argmax(cv_scores)]
print(f"Best k: {best_k}")

# Step 4: Plot Accuracy vs. k
plt.figure(figsize=(8, 5))
plt.plot(k_values, cv_scores, marker='o', linestyle='dashed',
color='b')
plt.xlabel('Number of Neighbors (k)')
plt.ylabel('Cross-Validated Accuracy')
plt.title('Finding the Best k for KNN')
plt.xticks(k_values)
plt.grid()
plt.show()

# Step 5: Train final model with the best k
best_knn_model = KNeighborsClassifier(n_neighbors=best_k)
best_knn_model.fit(X_train_resampled, y_train_resampled)

# Step 6: Make predictions on scaled test data
y_pred_best = best_knn_model.predict(X_test_scaled)

# Step 7: Evaluate the final model
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
accuracy_best = accuracy_score(y_test, y_pred_best)
print(f"Final Model Accuracy with k={best_k}: {accuracy_best *
100:.2f}%")

conf_matrix = confusion_matrix(y_test, y_pred_best)
class_report = classification_report(y_test, y_pred_best)

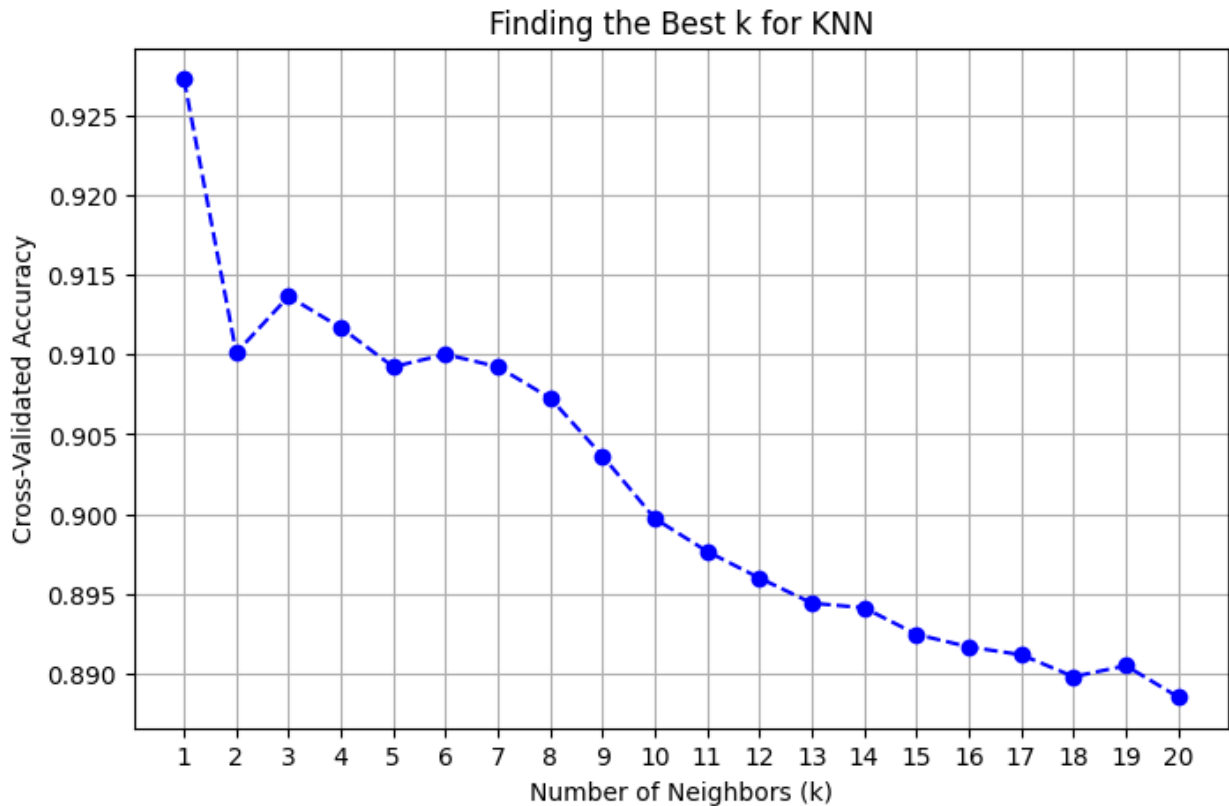
```

```

print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)

```

Best k: 1



Final Model Accuracy with k=1: 88.25%

Confusion Matrix:

```

[[1585  114]
 [ 124  203]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.93	0.93	1699
1	0.64	0.62	0.63	327
accuracy			0.88	2026
macro avg	0.78	0.78	0.78	2026
weighted avg	0.88	0.88	0.88	2026

```

import numpy as np
from sklearn.preprocessing import MinMaxScaler

```

```

# Load your trained model and scaler
# model = your_trained_model
# scaler = your_fitted_scaler

# Selected features as per your specification
selected_features = [
    'Total_Relationship_Count',
    'Total_Revolving_Bal',
    'Total_Ct_Chng_Q4_Q1',
    'Total_Trans_Ct',
    'Contacts_Count_12_mon',
    'Months_Inactive_12_mon'
]

# Creating a single sample using the values you provided
new_sample = np.array([[ -1.5          ,  0.3125          , -0.10212766,
 0.88888889,  1.          ],
                        [ 1.          ]])

# Scale the features using your trained scaler
# scaled_sample = scaler.transform(new_sample)

# Make prediction
# prediction = model.predict(scaled_sample)
# probability = model.predict_proba(scaled_sample)[0][1] #
# Probability of attrition (class 1)

# For demonstration (replace with actual prediction)
prediction = best_knn_model.predict(new_sample)

print("Customer data:", new_sample[0])
print("Prediction:", "Will attrite" if prediction == 1 else "Will
stay")

Customer data: [-1.5          0.3125          -0.10212766  0.88888889  1.
 1.          ]
Prediction: Will stay

X_train_resampled
array([[ -1.5          ,  0.3125          , -0.10212766,  0.88888889,  1.
 ,
         1.          ],
       [ -1.5          ,  0.36401099, -1.31489362, -0.77777778,  1.
 ,
         2.          ],
       [ 1.          ,  0.25755495,  2.33191489, -0.86111111,  1.
 ,
        -1.          ],
       ...,

```

```
[-1.      ,  0.8543956 , -0.26857441, -0.53821286,  1.
,
  1.10828913],
[-0.5      , -0.75829704, -0.3301146 , -1.10582922,  1.
,
  1.      ],
[-0.2969941 , -0.36722024, -1.59248204, -0.64411046, -1.
,
  1.      ]])
```